

Belinda Z. Li

CONTACT ✉ bzl@anthropic.com 🌐 belindal.github.io

ACADEMIC POSITIONS **University of Chicago** 08/2027 –
Assistant Professor

EDUCATION **Massachusetts Institute of Technology** 2020 – 2026
PhD, Electrical Engineering and Computer Science
- Advisor: Jacob Andreas
- Thesis: World Models, User Models, and Self Models in AI Systems

University of Washington 2015 – 2019
Bachelors of Science, Computer Science, Magna Cum Laude
- Advisor: Luke Zettlemoyer

EMPLOYMENT **Anthropic** 06/2026 –
Member of Technical Staff

Massachusetts Institute of Technology 09/2020 – 05/2026
Research Assistant, with Jacob Andreas

Transluce AI 05/2025 – 12/2025
Research Fellow

Letta AI 05/2025 – 12/2025
Research Fellow

Google Deepmind 06/2024 – 01/2025
Student Researcher

Microsoft Semantic Machines 06/2022 – 08/2022
Research Intern

Facebook AI 06/2021 – 08/2021
Research Intern, AI Integrity Team

Facebook AI 08/2019 – 09/2020
Software Engineer, AI Integrity Team (Facebook AI Applied Research)

Facebook 06/2018 – 08/2018
Software Engineer Intern, Bot Detection/Traffic Quality Team

PREPRINTS

[1] **Belinda Z. Li***, Itamar Pres*, Laura Ruis*, Zifan Carl Guo, Keya Hu, Mehul Damani, Isha Puri, Ekdeep Singh Lubana, and Jacob Andreas. *It's Time to Optimize for Self-Consistency*. Manuscript in submission. 2026.

[2] **Belinda Z. Li**, Zifan Carl Guo, Vincent Huang, Jacob Steinhardt, and Jacob Andreas. *Training Language Models to Explain Their Own Computations*. 2025.

[8] Kunal Handa, Yarin Gal, Ellie Pavlick, Noah Goodman, Jacob Andreas, Alex Tamkin, and **Belinda Z. Li**. *Bayesian Preference Elicitation with Language Models*. 2024.

- [22] Sinong Wang, **Belinda Z. Li**, Madian Khabsa, Han Fang, and Hao Ma. *Linformer: Self-Attention with Linear Complexity*. 2020.
- [23] Qinyuan Ye, **Belinda Z. Li**, Sinong Wang, Benjamin Bolte, Hao Ma, Xiang Ren, Wen-tau Yih, and Madian Khabsa. *Studying Strategically: Learning to Mask for Closed-book QA*. 2020.

PUBLICATIONS

- [3] **Belinda Z. Li**, Zifan Carl Guo, and Jacob Andreas. “(How) Do Language Models Track State?” In: *Forty-second International Conference on Machine Learning*. 2025. URL: <https://openreview.net/forum?id=8SXosAVIFH>.
- [4] **Belinda Z. Li**, Been Kim, and Zi Wang. “QuestBench: Can LLMs ask the right question to acquire information in reasoning tasks?” In: *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*. 2025. URL: <https://openreview.net/forum?id=gpwA9aZLTZ>.
- [5] **Belinda Z. Li***, Alex Tamkin*, Noah Goodman, and Jacob Andreas. “Eliciting Human Preferences with Language Models”. In: *The Thirteenth International Conference on Learning Representations*. 2025. URL: <https://openreview.net/forum?id=LvDwwAgMEW>.
- [6] **Belinda Z. Li**, Emmy Liu, Alexis Ross, Abbas Zeitoun, Graham Neubig, and Jacob Andreas. “Language Modeling with Editable External Knowledge”. In: *Findings of 2025 Annual Conference of the Nations of the Americas Chapter of the ACL*. 2025. URL: <https://openreview.net/forum?id=x5pPbgRR0t>.
- [7] Andi Peng, **Belinda Z. Li**, Ilia Sucholutsky, Nishanth Kumar, Julie Shah, Jacob Andreas, and Andreea Bobu. “Adaptive Language-Guided Abstraction from Contrastive Explanations”. In: *8th Annual Conference on Robot Learning*. 2024. URL: <https://openreview.net/forum?id=OGjGtN6hoo>.
- [9] Andi Peng, Andreea Bobu, **Belinda Z. Li**, Theodore Summers, Ilia Sucholutsky, Nishanth Kumar, Thomas Griffiths, and Julie Shah. “Preference-Conditioned Language-Guided Abstractions”. In: *Proceedings of the ACM/IEEE International Conference on Human Robot Interaction (HRI 2024)*. ACM/IEEE. 2024.
- [10] Andi Peng, Ilia Sucholutsky, **Belinda Z. Li**, Theodore Summers, Thomas Griffiths, Jacob Andreas, and Julie Shah. “Learning with language-guided state abstractions”. In: *RSS Workshop on Social Intelligence in Humans and Robots*. 2023.
- [11] **Belinda Z. Li**, Jason Eisner, Adam Pauls, and Sam Thomson. “*Toward Interactive Dictation*”. In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Toronto, Canada: Association for Computational Linguistics, July 2023, pp. 15319–15338.
- [12] Evan Hernandez, **Belinda Z. Li**, and Jacob Andreas. “Inspecting and Editing Knowledge Representations in Language Models”. In: *First Conference on Language Modeling*. 2024. URL: <https://openreview.net/forum?id=ADtL6fgNRv>.
- [13] **Belinda Z. Li**, William Chen, Pratyusha Sharma, and Jacob Andreas. “LaMPP: Language Models as Probabilistic Priors for Perception and Action”. In: *ICLR 2024 Workshop on Generative Models for Decision Making*. under review. 2024. URL: <https://openreview.net/forum?id=1aMfd925b5>.

- [14] **Belinda Z. Li**, Maxwell Nye, and Jacob Andreas. “[Language Modeling with Latent Situations](#)”. In: *Findings of the Association for Computational Linguistics: ACL 2023*. Toronto, Canada: Association for Computational Linguistics, July 2023, pp. 12556–12571.
- [15] **Belinda Z. Li**, Jane Yu, Madian Khabsa, Luke Zettlemoyer, Alon Halevy, and Jacob Andreas. “[Quantifying Adaptability in Pre-trained Language Models with 500 Tasks](#)”. In: *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Seattle, United States: Association for Computational Linguistics, July 2022, pp. 4696–4715.
- [16] Qinyuan Ye, **Belinda Z. Li**, Sinong Wang, Benjamin Bolte, Hao Ma, Wen-tau Yih, Xiang Ren, and Madian Khabsa. “[On the Influence of Masking Policies in Intermediate Pre-training](#)”. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Online and Punta Cana, Dominican Republic: Association for Computational Linguistics, Nov. 2021, pp. 7190–7202.
- [17] **Belinda Z. Li**, Maxwell Nye, and Jacob Andreas. “[Implicit Representations of Meaning in Neural Language Models](#)”. In: *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Online: Association for Computational Linguistics, Aug. 2021, pp. 1813–1827.
- [18] Nayeon Lee, **Belinda Z. Li**, Sinong Wang, Pascale Fung, Hao Ma, Wen-tau Yih, and Madian Khabsa. “[On Unifying Misinformation Detection](#)”. In: *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. Online: Association for Computational Linguistics, June 2021, pp. 5479–5485.
- [19] **Belinda Z. Li**, Sewon Min, Srinivasan Iyer, Yashar Mehdad, and Wen-tau Yih. “[Efficient One-Pass End-to-End Entity Linking for Questions](#)”. In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Online: Association for Computational Linguistics, Nov. 2020, pp. 6433–6441.
- [20] Nayeon Lee, **Belinda Z. Li**, Sinong Wang, Wen-tau Yih, Hao Ma, and Madian Khabsa. “[Language Models as Fact Checkers?](#)” In: *Proceedings of the Third Workshop on Fact Extraction and VERification (FEVER)*. Association for Computational Linguistics, July 2020, pp. 36–41.
- [21] **Belinda Z. Li**, Gabriel Stanovsky, and Luke Zettlemoyer. “[Active Learning for Coreference Resolution using Discrete Annotation](#)”. In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, July 2020, pp. 8320–8331.

AWARDS	Rising Stars in EECS	2024
	Clare Boothe Luce Graduate Fellowship	2021
	National Defence Science and Engineering Graduate Fellowship	2021
	National Science Foundation Graduate Research Fellowship (<i>Declined</i>)	2021
	Ida M. Green Memorial Fellowship	2020 – 2021
	Denice Dee Denton Scholars Endowment	2018 – 2019
	UW Dean’s List	2015 – 2017 (annual), SP18/F18/SP19 (quarterly)
	National Merit Scholar	2016
INVITED TALKS	[TBD]	
	Mechanistic Interpretability Workshop (MechInterp) @ ICML2026 (Lightning Talk)	07/2026

	[TBD]	
	Epistemic Intelligence in Machine Learning Workshop (EIML) @ ICML2026	07/2026
	<i>World Models, User Models, and Self Models in AI Systems</i>	
	Argonne National Laboratory Seminar	06/2026
	<i>Introspective Interpretability: Training Language Models to Explain Their Own Computations</i>	
	Oxford University	03/2026
	Far.AI Alignment Workshop (Lightning Talk)	03/2026
	Harvard University (Wattenberg + Viegas Lab)	01/2026
	TTIC Young Researcher Seminar Series	01/2026
	Northeastern University (Bau Lab)	01/2026
	<i>Solving the Specification Problem through Interaction</i>	
	Bacconi University (MilaNLP)	01/2026
	<i>QuestBench: Can LLMs Ask the Right Questions?</i>	
	Executive Code Podcast	07/2025
	<i>(How) Do Language Models Track State?</i>	
	Plutous AI Fireside Chat	04/2025
	<i>Eliciting Tasks and Preferences with Language Models</i>	
	Toronto Data Workshop	05/2024
	<i>World State in Language Models</i>	
	MILA	10/2023
	ServiceNow	10/2023
	<i>Quantifying Adaptability in Pre-trained Language Models with 500 Tasks</i>	
	NAACL Conference (Oral)	07/2022
	<i>Implicit Representations of Meaning in Neural Language Models</i>	
	Google N2Formal Team	09/2021
	MIT EI Workshop	09/2021
	Google DeepMind	07/2021
	Primer AI	06/2021
MENTORSHIP	PhDs	
	Itamar Pres	09/2025 –
	Masters	
	Zifan Carl Guo	09/2024 – 06/2026
	Kunal Handa (→ Anthropic MTS)	07/2023 – 03/2024
	Will Chen (→ UC Berkeley PhD)	05/2022 – 02/2023
	Undergraduates	
	George Cao	07/2025 –
	Amiri Hayes	06/2025 –
	Nolan Clement	09/2024 – 05/2025
	Maxine Liu (→ UPenn PhD)	06/2024 – 02/2025
	Nitya Parthasarathy	09/2020 – 06/2021
	Other	
	Tristan Day (CBAI)	03/2026 –
TEACHING	6.861*: Quantitative Methods for NLP	Fall 2023
	Teaching Assistant, Massachusetts Institute of Technology	

CSE 311: Foundations of Computing I (Discrete Math) <i>Teaching Assistant, University of Washington (Instructor: Kevin Zatloukal)</i>	Spring 2019
CSE 312: Foundations of Computing II (Probability/Statistics) <i>Teaching Assistant, University of Washington (Instructor: Martin Tompa)</i>	Winter 2019
CSE 312: Foundations of Computing II (Probability/Statistics) <i>Teaching Assistant, University of Washington (Instructor: Martin Tompa)</i>	Winter 2018
CSE 331: Software Design and Implementation <i>Teaching Assistant, University of Washington (Instructor: Kevin Zatloukal)</i>	Autumn 2017

SERVICE	Grant Reviewer - Schmidt Sciences AI Interpretability RFP, 2026 Area Chair - Mechanistic Interpretability Workshop @ ICML 2026 Program Committee / Reviewer - NeurIPS 2026 - ICLR 2024, 2025 - ICML 2024, 2025, 2026 - ACL 2023 - ARR February 2025; December, October, June, April, February 2024 - COLING 2022, 2020 - EMNLP 2022, 2021, 2020 Workshop Organizing - Assessing World Models: Methods and Metrics for Evaluating Understanding @ ICML 2025, (Co-Lead Organizer) - Wordplay: When Language Meets Games @ ACL 2024 MIT Summer Research Program Mentor MIT EECS Department Graduate Admissions Committee MIT GAAP Mentor Student Volunteer/Scholarship - ACL 2021, 2020	2024, 2025 2023 2020, 2021
---------	--	---

LAST
UPDATED

June 24, 2026