

Belinda Z. Li

CONTACT ✉ bzl@anthropic.com ✉ bzli@uchicago.edu 🌐 belindal.github.io

CURRENT POSITION **Anthropic** 06/2026 –
Member of Technical Staff

ACADEMIC POSITIONS **University of Chicago** 08/2027 –
Assistant Professor

EDUCATION **Massachusetts Institute of Technology** 2020 – 2026
PhD, Electrical Engineering and Computer Science
- Advisor: Jacob Andreas
- Thesis: World Models, User Models, and Self Models in AI Systems

University of Washington 2015 – 2019
Bachelor of Science, Computer Science, Magna Cum Laude
- Advisor: Luke Zettlemoyer

EMPLOYMENT **Massachusetts Institute of Technology** 09/2020 – 05/2026
Research Assistant, with Jacob Andreas

Transluce AI 05/2025 – 12/2025
Research Fellow

Letta AI 05/2025 – 12/2025
Research Fellow

Google DeepMind 06/2024 – 01/2025
Student Researcher

Microsoft Semantic Machines 06/2022 – 08/2022
Research Intern

Facebook AI 06/2021 – 08/2021
Research Intern, AI Integrity Team

Facebook AI 08/2019 – 09/2020
Software Engineer, AI Integrity Team (Facebook AI Applied Research)

Facebook 06/2018 – 08/2018
Software Engineer Intern, Bot Detection/Traffic Quality Team

PUBLICATIONS [1] Zifan Carl Guo, Laura Ruis, Jacob Andreas, and **Belinda Z. Li**. “[Introspective Coupling: Self-Explanation Training Tracks Behavioral Change Despite Fixed Supervision](#)”. In: *The Fortieth Annual Conference on Neural Information Processing Systems*. 2026.
[2] **Belinda Z. Li**, Zifan Carl Guo, Vincent Huang, Jacob Steinhardt, and Jacob Andreas. “[Training Language Models to Explain Their Own Computations](#)”. In: *The Fortieth Annual Conference on Neural Information Processing Systems*. 2026.

- [3] Amiri Hayes, **Belinda Z. Li**, and Jacob Andreas. “Explaining Attention with Program Synthesis”. In: *InterpScience Workshop at NeurIPS 2026*. 2026.
- [4] Itamar Pres, Laura Ruis, Melat Ghebreselassie, **Belinda Z. Li***, and Jacob Andreas*. “Self-CTRL: Self-Consistency Training with Reinforcement Learning”. In: *Actionable Interpretability Workshop at COLM 2026*. Oral presentation. 2026.
- [5] **Belinda Z. Li***, Itamar Pres*, Laura Ruis*, Zifan Carl Guo, Keya Hu, Mehul Damani, Isha Puri, Ekdeep Singh Lubana, and Jacob Andreas. “Position: It’s Time to Optimize for Self-Consistency”. In: *Forty-third International Conference on Machine Learning Position Paper Track*. 2026.
- [6] **Belinda Z. Li**, Zifan Carl Guo, and Jacob Andreas. “(How) Do Language Models Track State?”. In: *Forty-second International Conference on Machine Learning*. 2025.
- [7] **Belinda Z. Li**, Been Kim, and Zi Wang. “QuestBench: Can LLMs ask the right question to acquire information in reasoning tasks?”. In: *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*. 2025.
- [8] **Belinda Z. Li***, Alex Tamkin*, Noah Goodman, and Jacob Andreas. “Eliciting Human Preferences with Language Models”. In: *The Thirteenth International Conference on Learning Representations*. 2025.
- [9] **Belinda Z. Li**, Emmy Liu, Alexis Ross, Abbas Zeitoun, Graham Neubig, and Jacob Andreas. “Language Modeling with Editable External Knowledge”. In: *Findings of 2025 Annual Conference of the Nations of the Americas Chapter of the ACL*. 2025.
- [10] Andi Peng, **Belinda Z. Li**, Ilia Sucholutsky, Nishanth Kumar, Julie Shah, Jacob Andreas, and Andreea Bobu. “Adaptive Language-Guided Abstraction from Contrastive Explanations”. In: *8th Annual Conference on Robot Learning*. 2024.
- [11] Evan Hernandez, **Belinda Z. Li**, and Jacob Andreas. “Inspecting and Editing Knowledge Representations in Language Models”. In: *First Conference on Language Modeling*. 2024.
- [12] **Belinda Z. Li**, William Chen, Pratyusha Sharma, and Jacob Andreas. “LaMPP: Language Models as Probabilistic Priors for Perception and Action”. In: *ICLR 2024 Workshop on Generative Models for Decision Making*. 2024.
- [13] Andi Peng, Andreea Bobu, **Belinda Z. Li**, Theodore Sumers, Ilia Sucholutsky, Nishanth Kumar, Thomas Griffiths, and Julie Shah. “Preference-Conditioned Language-Guided Abstractions”. In: *Proceedings of the ACM/IEEE International Conference on Human Robot Interaction (HRI 2024)*. 2024.
- [14] **Belinda Z. Li**, Jason Eisner, Adam Pauls, and Sam Thomson. “Toward Interactive Dictation”. In: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2023.
- [15] **Belinda Z. Li**, Maxwell Nye, and Jacob Andreas. “Language Modeling with Latent Situations”. In: *Findings of the Association for Computational Linguistics: ACL 2023*. 2023.
- [16] Andi Peng, Ilia Sucholutsky, **Belinda Z. Li**, Theodore Sumers, Thomas Griffiths, Jacob Andreas, and Julie Shah. “Learning with Language-Guided State Abstractions”. In: *RSS Workshop on Social Intelligence in Humans and Robots*. 2023.

- [17] **Belinda Z. Li**, Jane Yu, Madian Khabsa, Luke Zettlemoyer, Alon Halevy, and Jacob Andreas. “[Quantifying Adaptability in Pre-trained Language Models with 500 Tasks](#)”. In: *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. 2022.
- [18] Qinyuan Ye, **Belinda Z. Li**, Sinong Wang, Benjamin Bolte, Hao Ma, Wen-tau Yih, Xiang Ren, and Madian Khabsa. “[On the Influence of Masking Policies in Intermediate Pre-training](#)”. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 2021.
- [19] **Belinda Z. Li**, Maxwell Nye, and Jacob Andreas. “[Implicit Representations of Meaning in Neural Language Models](#)”. In: *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. 2021.
- [20] Nayeon Lee, **Belinda Z. Li**, Sinong Wang, Pascale Fung, Hao Ma, Wen-tau Yih, and Madian Khabsa. “[On Unifying Misinformation Detection](#)”. In: *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. 2021.
- [21] **Belinda Z. Li**, Sewon Min, Srinivasan Iyer, Yashar Mehdad, and Wen-tau Yih. “[Efficient One-Pass End-to-End Entity Linking for Questions](#)”. In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. 2020.
- [22] Nayeon Lee, **Belinda Z. Li**, Sinong Wang, Wen-tau Yih, Hao Ma, and Madian Khabsa. “[Language Models as Fact Checkers?](#)” In: *Proceedings of the Third Workshop on Fact Extraction and VERification (FEVER)*. 2020.
- [23] **Belinda Z. Li**, Gabriel Stanovsky, and Luke Zettlemoyer. “[Active Learning for Coreference Resolution using Discrete Annotation](#)”. In: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*. 2020.

PREPRINTS

- [24] Kunal Handa, Yarin Gal, Ellie Pavlick, Noah Goodman, Jacob Andreas, Alex Tamkin, and **Belinda Z. Li**. “[Bayesian Preference Elicitation with Language Models](#)”. 2024.
- [25] Sinong Wang, **Belinda Z. Li**, Madian Khabsa, Han Fang, and Hao Ma. “[Linformer: Self-Attention with Linear Complexity](#)”. 2020.
- [26] Qinyuan Ye, **Belinda Z. Li**, Sinong Wang, Benjamin Bolte, Hao Ma, Xiang Ren, Wen-tau Yih, and Madian Khabsa. “[Studying Strategically: Learning to Mask for Closed-book QA](#)”. 2020.

AWARDS

Rising Stars in EECS	2024
Clare Boothe Luce Graduate Fellowship	2021
National Defense Science and Engineering Graduate Fellowship	2021
National Science Foundation Graduate Research Fellowship (<i>Declined</i>)	2021
Ida M. Green Memorial Fellowship	2020 – 2021
Denice Dee Denton Scholars Endowment	2018 – 2019
UW Dean’s List	2015 – 2017 (annual), SP18/F18/SP19 (quarterly)
National Merit Scholar	2016

INVITED TALKS

Actionable Interpretability Workshop @ COLM 2026 (Panelist)	10/2026
<i>Consistency is an Interpretability Pre-requisite</i>	
New England Mechanistic Interpretability Workshop (NEMI)	08/2026

<i>AI Systems Should Have Accurate Self-Models</i>	
Mechanistic Interpretability Workshop (MechInterp) @ ICML 2026 (Lightning Talk)	07/2026
<i>Resolving Uncertainty via Interaction</i>	
Epistemic Intelligence in Machine Learning Workshop (EIML) @ ICML 2026	07/2026
<i>World Models, User Models, and Self Models in AI Systems</i>	
Argonne National Laboratory Seminar	06/2026
<i>Introspective Interpretability: Training Language Models to Explain Their Own Computations</i>	
Oxford University	03/2026
Far.AI Alignment Workshop (Lightning Talk)	03/2026
Harvard University (Wattenberg + Viegas Lab)	01/2026
TTIC Young Researcher Seminar Series	01/2026
Northeastern University (Bau Lab)	01/2026
<i>Solving the Specification Problem through Interaction</i>	
Bocconi University (MilaNLP)	01/2026
<i>QuestBench: Can LLMs Ask the Right Questions?</i>	
Executive Code Podcast	07/2025
<i>(How) Do Language Models Track State?</i>	
Plutous AI Fireside Chat	04/2025
<i>Eliciting Tasks and Preferences with Language Models</i>	
Toronto Data Workshop	05/2024
<i>World State in Language Models</i>	
MILA	10/2023
ServiceNow	10/2023
<i>Quantifying Adaptability in Pre-trained Language Models with 500 Tasks</i>	
NAACL Conference (Oral)	07/2022
<i>Implicit Representations of Meaning in Neural Language Models</i>	
Google N2Formal Team	09/2021
MIT EI Workshop	09/2021
Google DeepMind	07/2021
Primer AI	06/2021

MENTORSHIP

PhDs

Itamar Pres 09/2025 – 06/2026

Masters

Zifan Carl Guo 09/2024 –

Kunal Handa (→ Anthropic MTS) 07/2023 – 03/2024

Will Chen (→ UC Berkeley PhD) 05/2022 – 02/2023

Undergraduates

Melat Ghebreselassie 02/2026 – 06/2026

George Cao 07/2025 –

Amiri Hayes 06/2025 –

Nolan Clement 09/2024 – 05/2025

Maxine Liu (→ UPenn PhD) 06/2024 – 02/2025

Nitya Parthasarathy 09/2020 – 06/2021

	Other Arav Dhoot (SPAR) 09/2026 – Bryan Maruyama (SPAR) 09/2026 – Laura Gomezjurado (SPAR) 09/2026 – Tobias Moser (SPAR) 09/2026 – Tommaso Cerruti (SPAR) 09/2026 – Victoria Li (Anthropic Fellows) 08/2026 – Yuxin Xie (Anthropic Fellows) 08/2026 – Jiawei Li (Anthropic Fellows) 08/2026 – Tristan Day (CBAI) 03/2026 –
TEACHING	6.861*: Quantitative Methods for NLP Fall 2023 <i>Teaching Assistant, Massachusetts Institute of Technology</i> CSE 311: Foundations of Computing I (Discrete Math) Spring 2019 <i>Teaching Assistant, University of Washington (Instructor: Kevin Zatloukal)</i> CSE 312: Foundations of Computing II (Probability/Statistics) Winter 2019 <i>Teaching Assistant, University of Washington (Instructor: Martin Tompa)</i> CSE 312: Foundations of Computing II (Probability/Statistics) Winter 2018 <i>Teaching Assistant, University of Washington (Instructor: Martin Tompa)</i> CSE 331: Software Design and Implementation Autumn 2017 <i>Teaching Assistant, University of Washington (Instructor: Kevin Zatloukal)</i>
SERVICE	Grant Reviewer - Schmidt Sciences AI Interpretability RFP, 2026 Area Chair - Mechanistic Interpretability Workshop @ ICML 2026 Program Committee / Reviewer - NeurIPS 2026 - ICLR 2024, 2025 - ICML 2024, 2025, 2026 - ACL 2023 - ARR February 2025; December, October, June, April, February 2024 - COLING 2022, 2020 - EMNLP 2022, 2021, 2020 Workshop Organizing - Assessing World Models: Methods and Metrics for Evaluating Understanding @ ICML 2025 (<i>Co-Lead Organizer</i>) - Wordplay: When Language Meets Games @ ACL 2024 MIT Summer Research Program Mentor 2024, 2025 MIT EECS Department Graduate Admissions Committee 2023 MIT GAAP Mentor 2020, 2021 Student Volunteer/Scholarship - ACL 2021, 2020
LAST UPDATED	October 8, 2026